

Airbnb Price Prediction Analysis

Question

How well can the price (log price) of Airbnb rentals be predicted using machine learning models?

Process

- Identify columns with missing or NaN values
- Plan how to deal with each individual column
- Clean data
- Transform categorical variables into encoded features
- Train a base XGBoost model on the cleaned data
- Evaluate results
- Tune hyperparameters
- Evaluate model
- Plot feature importances

Problem areas:

There are a few problem area areas with this analysis.

The first one is that certain categorical features (neighbourhood), if entirely used and encoded, would lead to extreme levels of dimensionality. These features were not included as a result, with the idea that latitude and longitude may make up for them

The other main problem was processing power available. The machine used for this analysis would have taken days to do a comprehensive grid search of hyperparameters, so performing it in chunks was used instead. Unfortunately, this means that the hyperparameters are not perfectly tuned.

MAE - 0.272035938048597

MAPE - 0.05719580023745668

```
In [40]: #imports
import pandas as pd
import numpy as np
import seaborn as sns
```

```
import matplotlib.pyplot as plt
import plotly.graph_objects as px
import xgboost
import warnings
from xgboost import plot_tree

warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', 40)
```

```
In [41]: #read in data

data = pd.read_csv('../data/train.csv')

data.head()
```

```
Out[41]:
```

	id	log_price	property_type	room_type	amenities	accommodates	ba
0	6901257	5.010635	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitch...	3	
1	6304928	5.129899	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitch...	7	
2	7919400	4.976734	Apartment	Entire home/apt	{TV,"Cable TV","Wireless Internet","Air condit...	5	
3	13418779	6.620073	House	Entire home/apt	{TV,"Cable TV",Internet,"Wireless Internet",Ki...	4	
4	3808709	4.744932	Apartment	Entire home/apt	{TV,Internet,"Wireless Internet","Air conditio...	2	

```
In [42]: data.info()

#has missing values:
#bathrooms, first_review, host_has_profile_pic, host_identity_verified, host_respon
#neighbourhood, review_scores_rating, thumbnail_url, zipcode, bedrooms, beds
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74111 entries, 0 to 74110
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     74111 non-null  int64
1   log_price                             74111 non-null  float64
2   property_type                         74111 non-null  object
3   room_type                             74111 non-null  object
4   amenities                             74111 non-null  object
5   accommodates                          74111 non-null  int64
6   bathrooms                             73911 non-null  float64
7   bed_type                              74111 non-null  object
8   cancellation_policy                   74111 non-null  object
9   cleaning_fee                          74111 non-null  bool
10  city                                   74111 non-null  object
11  description                            74111 non-null  object
12  first_review                           58247 non-null  object
13  host_has_profile_pic                   73923 non-null  object
14  host_identity_verified                 73923 non-null  object
15  host_response_rate                     55812 non-null  object
16  host_since                             73923 non-null  object
17  instant_bookable                       74111 non-null  object
18  last_review                           58284 non-null  object
19  latitude                               74111 non-null  float64
20  longitude                              74111 non-null  float64
21  name                                   74111 non-null  object
22  neighbourhood                           67239 non-null  object
23  number_of_reviews                      74111 non-null  int64
24  review_scores_rating                   57389 non-null  float64
25  thumbnail_url                          65895 non-null  object
26  zipcode                                73145 non-null  object
27  bedrooms                               74020 non-null  float64
28  beds                                   73980 non-null  float64
dtypes: bool(1), float64(7), int64(3), object(18)
memory usage: 15.9+ MB

```

```

In [43]: from sklearn.impute import SimpleImputer
         from sklearn.model_selection import train_test_split

#cleaning/transform:

#date to day, month, year cols: first_review, host_since, last_review
#remove: zipcode, description, name, id
#onehot encode (remove first): property_type, room_type, amenities, bed_type, cancellations
#bool to 1,0: cleaning_fee, thumbnail_url, instant_bookable

#has missing values:

#float (impute): bathrooms, review_scores_rating, bedrooms, beds
#date (-1): first_review, host_since, last_review
#str (fill "unknown"): neighbourhood
#str to bool (nan to 0): host_has_profile_pic, host_identity_verified
#str to int (remove percentages, impute): host_response_rate
#to bool (has value or not): thumbnail_url

```

```

data['first_review_day'], data['first_review_year'], data['first_review_month'] = p
data['host_since_day'], data['host_since_year'], data['host_since_month'] = pd.to_d
data['last_review_day'], data['last_review_year'], data['last_review_month'] = pd.t
data = data.drop(['first_review', 'host_since', 'last_review', 'zipcode', 'descript

categorical_variables = ['property_type', 'room_type', 'bed_type', 'cancellation_pol
amenities = data['amenities']
data['neighbourhood'] = data['neighbourhood'].replace(np.nan, 'Unknown')
data = data.drop(['amenities', 'neighbourhood'], axis=1) #neighborhood raises mae
data['cleaning_fee'] = data['cleaning_fee'].map({True: 1, False: 0})
data['host_has_profile_pic'] = data['host_has_profile_pic'].map({'t': 1, 'f': 0, np
data['host_identity_verified'] = data['host_identity_verified'].map({'t': 1, 'f': 0
data['instant_bookable'] = data['instant_bookable'].map({'t': 1, 'f': 0, np.nan: 0})
data['host_response_rate'] = data['host_response_rate'].astype(str).apply(lambda x:
data['thumbnail_url'] = data['thumbnail_url'].map(lambda x: 0 if pd.isna(x) else 1)
data['host_response_rate'] = data['host_response_rate'].astype('Float64')

date_cols = ['first_review_day', 'first_review_year', 'first_review_month', 'host_s
data[date_cols] = data[date_cols].apply(lambda x: x.fillna(-1))

imputed_cols = ['bedrooms', 'bathrooms', 'beds', 'review_scores_rating', 'host_resp
imputer = SimpleImputer(strategy='mean')
imputer.fit(data[imputed_cols])
data[imputed_cols] = imputer.transform(data[imputed_cols])

```


```

In [44]: #plot out different property types

fig, ax = plt.subplots(figsize=(12,12))
fig = sns.barplot(x='log_price', data=data.groupby('property_type', as_index=False)
fig.bar_label(fig.containers[0], padding=4)
fig.set_xlabel('Count')
fig.set_ylabel('Property Type')
fig.set_title('Counts of Property Types')

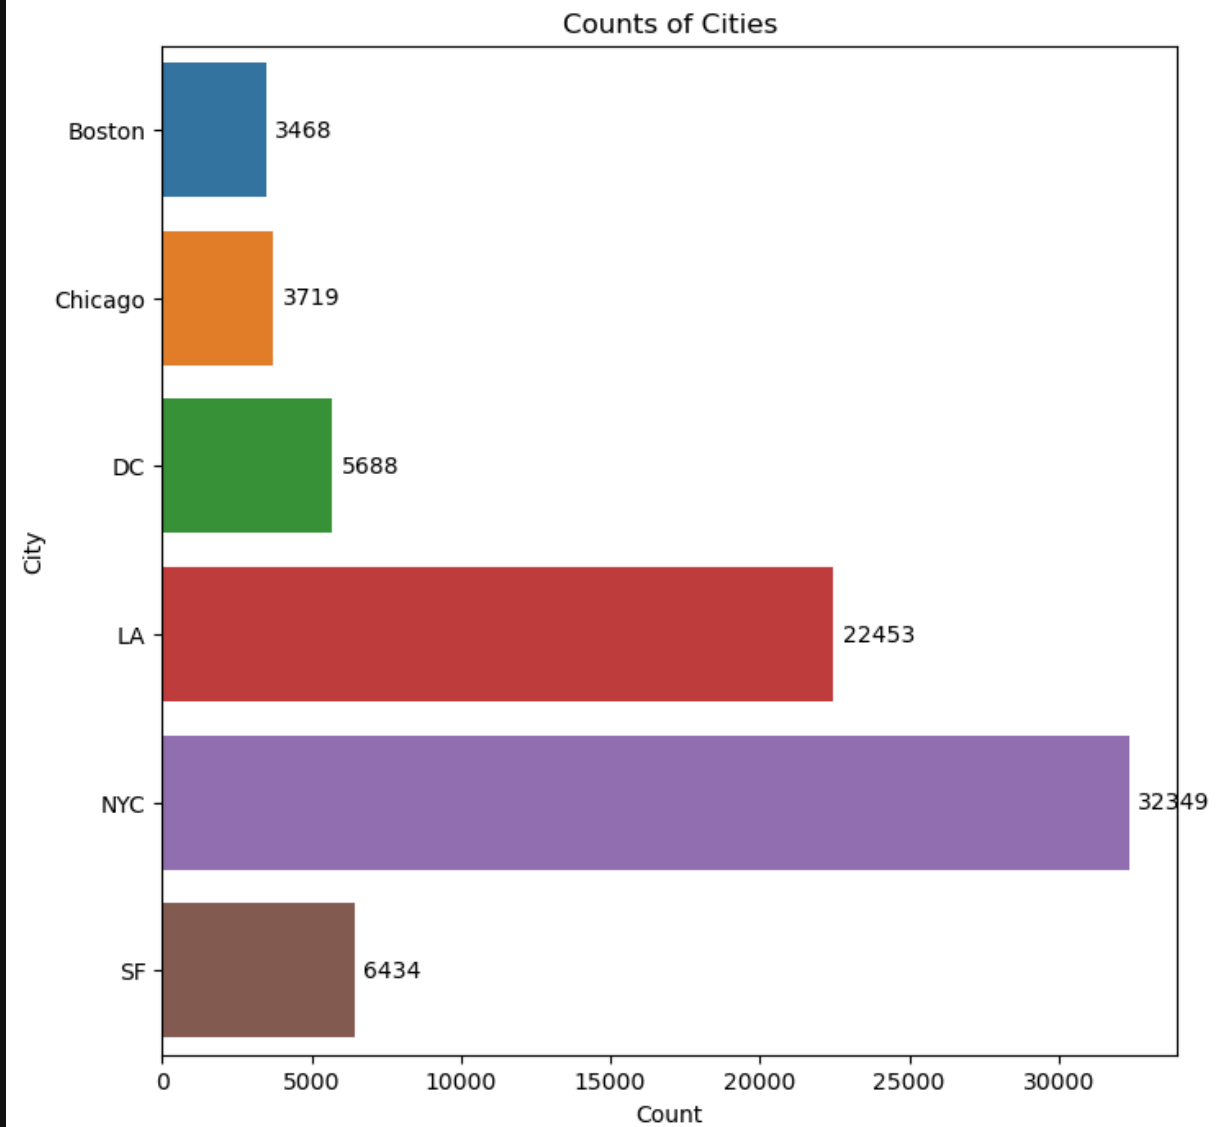
```

Out[44]: Text(0.5, 1.0, 'Counts of Property Types')



```
In [45]: #plot out city counts
fig, ax = plt.subplots(figsize=(8,8))
fig = sns.barplot(x='log_price', data=data.groupby('city', as_index=False).count()[
fig.bar_label(fig.containers[0], padding=4)
fig.set_xlabel('Count')
fig.set_ylabel('City')
fig.set_title('Counts of Cities')
```

```
Out[45]: Text(0.5, 1.0, 'Counts of Cities')
```



```
In [46]: #encode categorical variables, train/test split

data = pd.get_dummies(data, columns=categorical_variables, drop_first=True)
X = data.drop('log_price', axis=1)
y = data['log_price']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=8)
```

```
In [47]: #find shape after categorical variables are encoded

data.shape
```

```
Out[47]: (74111, 73)
```

```
In [48]: #initialize xgboost and fit to training data

from xgboost import XGBRegressor

boost = XGBRegressor()
```

```
boost.fit(X_train, y_train)
boost
```

Out[48]:

	XGBRegressor
	XGBRegressor(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=None, device=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=None, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=None, max_leaves=None, min_child_weight=None, missing=None, monotone_constraints=None, multi_output_label_func=None, n_estimators=None, num_parallel_tree=None, print_eval_information=None, random_state=None, raw_score=False, restore_from=None, scale_pos_weight=None, score_iteration=None, subsample=None, tree_method=None, validate_parameters=None, verbosity=None, watchlog=None, weight_postfix=None)

In [49]:

```
#model predictions
y_pred = boost.predict(X_test)
```

In [50]:

```
#model MAPE and MAE

from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error

print(f'MAPE: {mean_absolute_percentage_error(y_pred, y_test)}, MAE: {mean_absolute_error(y_pred, y_test)})

MAPE: 0.05878560869465756, MAE: 0.2794164988969564
```

In [23]:

[illegible]

```

        subsample=0.8,
        colsample_bytree=0.6,
        seed=8,
    ),
    param_grid=params, n_jobs=4, scoring='neg_mean_absolute_

grid_search.fit(X_train, y_train)
print(grid_search.best_params_, grid_search.best_score_)

```

```

{'booster': 'gbtree', 'eval_metric': 'rmse', 'learning_rate': 0.05, 'objective': 're
g:squarederror', 'random_state': 8, 'reg_alpha': 0.1, 'reg_lambda': 0} -0.2745637143
0471654

```

In [39]: *#tuned model MAPE and MAE*

```

print(f'MAPE: {mean_absolute_percentage_error(y_pred, y_test)}, MAE: {mean_absolute

MAPE: 0.05719580023745668, MAE: 0.272035938048597

```

In [35]: *#save graph as image (extremely large image)*

```

img = xgboost.to_graphviz(boost)
img.render('graph', format='png')

```

dot: graph is too large for cairo-renderer bitmaps. Scaling by 0.660732 to fit

(process:5328): GLib-GIO-WARNING **: 17:42:19.913: Unexpectedly, UWP app `18184where where.AndroidAppInstaller_0.1.25.0_x64__4v4sx105x6y4r' (AUMId `18184wherewhere.Andro idAppInstaller_4v4sx105x6y4r!App') supports 4 extensions but has no verbs

Out[35]: 'graph.png'

In [36]: *#plot feature importances*


```

feat_importances = pd.DataFrame(boost.get_booster().get_fscore().items(), columns=[

fig, ax = plt.subplots(figsize=(20,20))
sns.barplot(data=feat_importances.iloc[:len(feat_importances)//2], x='Importance',
ax.set_title('Feature Importances (1-32)')

```

Out[36]: Text(0.5, 1.0, 'Feature Importances (1-32)')



```
In [38]: #feature importances (cont.)
#NOTE, THESE ARE SMALLER IN IMPORTANCE THAN ABOVE GRAPH DESPITE BAR SIZE

fig, ax = plt.subplots(figsize=(20,20))
sns.barplot(data=feat_importances.iloc[len(feat_importances)//2:], x='Importance',
ax.set_title('Feature Importances (33-64)')
```

```
Out[38]: Text(0.5, 1.0, 'Feature Importances (33-64)')
```

